

MapReduce Programming

1. Enter the following URL in Google Chrome
<https://classroom-vc.gre.ac.uk>

Launch the Hadoop VM image, following the instructions from last week's lab.

2. Once the Hadoop desktop has loaded, right click on the desktop and start a terminal
3. At the command prompt type:

```
start-dfs.sh
start-yarn.sh
```

to start Hadoop.

4. Change directory into your workspace directory (if you don't have a workspace directory then starting Eclipse will create one for you)

```
cd ../workspace
```

5. Copy the Lab6 directory from Datasets into your workspace directory
6. Change current directory to the workspace/Lab6 directory
7. Open the WordCount.java file using Geany or Eclipse

```
import java.io.IOException;
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;

public class WordCount {

    public static class Map
        extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, OutputCollector<Text,
            IntWritable> output, Reporter reporter) throws IOException {

            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                output.collect(word, one);
            }
        }
    }

    public static class Reduce extends MapReduceBase
        implements Reducer<Text, IntWritable, Text, IntWritable> {

        public void reduce(Text key, Iterator<IntWritable> values,
            OutputCollector<Text, IntWritable> output, Reporter reporter) throws
            IOException {

```

```

        int sum = 0;
        while (values.hasNext()){
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}

public static void main(String[] args) throws Exception {
    JobConf conf = new JobConf(WordCount.class);
    conf.setJobName("wordcount");

    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);

    conf.setMapperClass(Map.class);
    conf.setCombinerClass(Reduce.class);
    conf.setReducerClass(Reduce.class);

    conf.setInputFormat(TextInputFormat.class);
    conf.setOutputFormat(TextOutputFormat.class);

    FileInputFormat.setInputPaths(conf, new Path(args[0]));
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));

    JobClient.runJob(conf);
}
}

```

8. Create a directory for the compiled WordCount classes

```
mkdir wc_classes
```

9. MapReduce programming in Java uses a number of libraries. Compile the WordCount.java file using libraries from the Hadoop Java API as follows:

```
javac -classpath $HADOOP_HOME/share/hadoop/common/hadoop-common-2.4.1.jar:$HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-client-core-2.4.1.jar:$HADOOP_HOME/share/hadoop/common/lib/hadoop-annotations-2.4.1.jar -d wc_classes WordCount.java
```

You can copy this command line from the java_command_line.txt file contained in the Lab6 folder

10. Package the compiled classes into a jar file

```
jar -cvf wordcount.jar -C wc_classes/ .
```

11. Put the test.txt file from last week's Lab onto the HDFS if it is not already there

```
hdfs dfs -put test.txt
```

12. Run the word count map reduce job using following in the command line:

```
hadoop jar wordcount.jar WordCount test.txt output
```

13. Check the results of the Word Count MapReduce, by copying the output directory from the HDFS to your local file system

```
hdfs dfs -get output
```

14. Try running Word Count on some of the books from Lab 5 (last week's lab).

15. If you are familiar with Java you might like to try modifying the WordCount.java example so that it counts hash tags:

- Open WordCount.java in Eclipse (or Geany) and save it to a new file called HashTagCount.java
- You will need to change the class name to HashTagCount and also the change parts of the main(...) method as follows:

```
JobConf conf = new JobConf(HashTagCount.class);  
conf.setJobName("hashtagcount");
```

- You will need to change the map(...) function to count occurrences of particular hashtags. You can use the startsWith() method of the java String class to detect hashtags. E.g.

```
if(token.startsWith("#")){...}
```

- The Reduce class can be left as it is
- You can compile, package and run your code using similar commands to those used for the WordCount example above.
- There is a data file called hashtags.txt in the Lab6 folder for you to test your code on.