

# Graph & Modern Databases

COMP1835

1

1

## Redis



2

2

## Objectives

- ▶ To introduce Redis data store
- ▶ To discuss CRUD commands
- ▶ To introduce different data types
- ▶ To learn to work with complex data types
- ▶ To discuss transactions
- ▶ To discuss the use of expiration

3

3

## Part 1

4

4

## Overview



- ▶ Redis is an open source, advanced **key-value store** and a solution for building high performance, scalable web applications.
- ▶ Redis stands for **RE**remote **D**ictionary **S**erver.
- ▶ Redis has three main features that sets it apart:
  - Redis holds its database entirely in the memory, using the disk only for persistence.
  - Redis has a relatively rich set of data types when compared to many key-value data stores.
  - Redis can replicate data to any number of slaves.
- ▶ 'Redis is like grease: it's most often used to lubricate moving parts and keep them working smoothly by reducing friction and speeding up their overall function.'

5

5

## Facts about Redis

- ▶ **Official Online Resources:**
  - <http://redis.io/>.
- ▶ **History:**
  - Project started in 2009 by Salvatore Sanfilippo. Salvatore created it for his startup company LLOOGG (<http://lloogg.com/>). Though still an independent project, Redis primary author is employed by VMware, who sponsors its development.
- ▶ **Technologies and Language:**
  - Implemented in C.
- ▶ **Access Methods:**
  - Rich set of methods and operations. Can access via Redis command-line interface and a set of well-maintained client libraries for languages like Java, Python, Ruby, C, C++, Lua, Haskell, AS3, and more.
- ▶ **Who Uses It:**
  - Twitter, Github, Stackoverflow, Pinterest, Snapchat, Craigslist, Flickr, etc.
- ▶ In 2015 Redis has been ranked the #4 NoSQL database in user satisfaction and market presence based on user reviews.

6

6

## Advantages



- ▶ Redis is exceptionally fast:
  - It can perform about 110,000 SETs per second, about 81,000 GETs per second.
- ▶ Redis supports rich data types:
  - It natively supports most of the datatypes such as list, set, sorted set, and hashes.
    - This makes it easy to solve a variety of problems since we know which problem can be handled better by which data type.
- ▶ Operations are atomic:
  - All Redis operations are atomic, which ensures that if two clients concurrently access, Redis server will receive the updated value.
- ▶ Redis is a multi-utility tool:
  - It can be used in a number of use cases such as caching, messaging-queues
  - Redis natively supports Publish/Subscribe, any short-lived data in your application, such as web application sessions, web page hit counts, etc.

7

7

## More ..

- ▶ Redis is not just a key-value store, but a different evolution path in the key-value DBs:
  - It supports more complex data types, though not to the degree that document-oriented database would
- ▶ Redis is an in-memory database but persistent on disk database
  - It represents a different trade off where very high write and read speed is achieved with the limitation of data sets that can't be larger than the memory.
  - To enhance the speed Redis purposely compromises the durability
    - In Redis, in the event of system failure or crash, Redis writes to disk but may fall behind and lose the data which is not stored.
- ▶ Redis supports set-based query operations but not with the granularity or type support you would find in a relational database.
- ▶ Redis is more of a toolkit of useful data structures algorithms and processes than a member of any specific database genre.

8

8

## Data Model

- ▶ Redis is a key-value store
- ▶ While the key must be a *String* information, the value could be one of these five different types: **STRINGS**, **LISTs**, **SETs**, **HASHes**, and **ZSETs**.
- ▶ An important difference between Redis and other key-value storage systems is that Redis supports not only strings, but also abstract data types
- ▶ Redis does not support:
  - automatic key allocation.
  - composite keys.
  - secondary indexes:

Key	Value
Name	Joe
Age	42
Occupation	Developer
WebPage	www.joe.co.uk

9

9

## Redis Commands

- ▶ Redis commands are used to perform some operations on Redis server.
- ▶ To run commands on Redis server, you need a Redis client, which is usually installed with Redis package
- ▶ To start Redis client, use **\$redis-cli**, which will connect to your local server.
- ▶ After you connect to Redis server running on the local machine, you can execute a command **PING**, that checks whether the server is running or not.
  - If you cannot connect, you'll receive an error message
  - Typing **help** will display a list of help options
- ▶ Example:

```
$redis-cli
redis 127.0.0.1:6379>
redis 127.0.0.1:6379> PING
PONG
redis 127.0.0.1:6379> help
```

10

10

## CRUD commands

► Syntax: `127.0.0.1:6379> COMMAND KEY_NAME`

► SET (Setting a Key):

```
127.0.0.1:6379> SET greeting "Hello World"
```

► GET (Getting a Key): `127.0.0.1:6379> GET greeting`  
"Hello World"

► DEL (Deleting a Key)

```
127.0.0.1:6379> GET greeting
"Hello World" // getting a key
127.0.0.1:6379> DEL greeting
(integer) 1 // key just got deleted
127.0.0.1:6379> GET greeting
(nil) // since key is deleted therefore, result is nil.
```

11

11

## Questions



- Name two advantages of using Redis data store.
- What are the limitations of Redis data store?
- Does Redis provide both **speed and durability**?

12

12

## Part 2

13

13

## Data Types – strings

### ► Strings

- Redis string is a sequence of bytes, you can store anything up to 512 megabytes in one string.

### ► String commands:

- **SET key value** - sets the value at the specified key.
- **GET key** - gets the value of a key.
- **GETRANGE key start end** - gets a substring of the string stored at a key.
- **GETSET key value** - sets the string value of a key and return its old value.
- **MGET key1 [key2..]** -gets the values of all the given keys as an ordered list
- **MSET key value [key value ...]** - sets multiple keys to multiple values
- **APPEND key value** - appends a value to a key
- **SETEX key seconds value** - Sets the value with the expiry of a key
- Full list of string command see here: <https://redis.io/commands#string>

### ► Example:

```
127.0.0.1:6379> SET greeting "Hello World"
127.0.0.1:6379> GET greeting
"Hello World" // getting a key
```

14

14

## Data Types – integers

- ▶ Although Redis stores strings, it recognizes integers and provides some simple operations for them.
- ▶ For example, if we want to keep a running total of how many key/value pairs are in our dataset, we can create a count and then increment it with the **INCR** command.
- ▶ Example:


```
127.0.0.1:6379> SET count 2
OK
127.0.0.1:6379> INCR count
(integer) 3
127.0.0.1:6379> GET count
"3"
```

- ▶ Although **GET** returns count as a string, **INCR** recognized it as an integer and added one to it.
- ▶ If the value can't be resolved to an integer, Redis will return an error.
  - You can also increment by any integer (**INCRBY**) or decrement (**DECR**, **DECRBY**).

15

15

## Complex Data Types



- ▶ Redis supports many complex data types, storing lists, hashes, sets, and sorted sets natively, which make it very attractive data store
- ▶ Different complex data types have appropriate commands associate with them, which generally follow a good pattern.
  - Set's commands begin with **S**,
  - Hashes' commands begin with **H**,
  - Sorted set's commands begin with **Z**.
  - List's commands generally start with either an **L** (for left) or an **R** (for right), depending on the direction of the operation (such as **LPUSH**).

16

16



## Data Types – Hashes

### Hashes

- A Redis hash is a collection of key value pairs. Redis Hashes are maps between string fields and string values. Hence, they are used to represent objects

### Hash commands:

- **HGET key field** - Gets the value of a hash field stored at the specified key.
- **HGETALL key** - Gets all the fields and values stored in a hash at the specified key
- **HVALS key** - Gets all the values in a hash
- **HLEN key** - Gets the number of fields in a hash
- **HMGET key field1 [field2]** - Gets the values of all the given hash fields
- **HSET key field value** - Sets the string value of a hash field
- **HMSET key field1 value1 [field2 value2 ]** - Sets multiple hash fields to multiple values
- Full list of hash command see here: <https://redis.io/commands#hash>

### Example:

```
127.0.0.1:6379> HMSET user:1 username nosql password nosqlpass
credits 15
OK
127.0.0.1:6379> HGETALL user:1
1) "username"
2) "nosql"
3) "password"
4) "nosqlpass"
5) "credits"
6) "15"
```

17

## Data Types –Lists

### Lists

- Redis Lists are simply lists of strings, sorted by insertion order. You can add elements to a Redis List on the head (L) or on the tail (R).

### Lists commands:

- **LINDEX key index** - Gets an element from a list by its index
- **LLEN key** - Gets the length of a list
- **LPUSH key value1 [value2]** - Adds one or multiple values to the beginning of the list
- **RPUSH key value1 [value2]** - Appends one or multiple values to the end of the list
- **LRANGE key start stop** - Gets a range of elements from a list
- Full list of list commands see here <https://redis.io/commands#list>
- All list operations in Redis use a zero based index.

### Example:

```
127.0.0.1:6379> LPUSH tutorials redis
(integer) 1
127.0.0.1:6379> LPUSH tutorials mongodb
(integer) 2
127.0.0.1:6379> LPUSH tutorials mysql
(integer) 3
127.0.0.1:6379> LRANGE tutorials 0 10
1) "mysql"
2) "mongodb"
3) "redis"
```

18

18

## Data Types – Sets

### ▶ Sets

- Redis Sets are **an unordered** collection of **unique** strings.
  - Unique means sets does not allow repetition of data in a key.
- Sets are an excellent choice for performing complex operations between two or more key values, such as unions or intersections.

### ▶ Sets commands:

- SADD key member1 [member2]** - Adds one or more members to a set
- SCARD key** - Gets the number of members in a set
- SMEMBERS key** - Gets all the members in a set
- Full list sets commands see here: <https://redis.io/commands#set>

### ◦ Example:

```
127.0.0.1:6379> SADD tutorials redis
(integer) 1
127.0.0.1:6379> SADD tutorials mongodb
(integer) 1
127.0.0.1:6379> SADD tutorials mysql
(integer) 1
127.0.0.1:6379> SADD tutorials mysql
(integer) 0
127.0.0.1:6379> SMEMBERS tutorials
1) "mysql"
2) "mongodb"
3) "redis"
```

19

## Operations with SETS

### ▶ Create two sets

```
127.0.0.1:6379> SADD fruit1 apples oranges
(integer) 2
127.0.0.1:6379> SADD fruit2 apples pears
(integer) 2
```

### ▶ To find the intersection of sets – values that are in both, we use the **SINTER** command.

```
127.0.0.1:6379> SINTER fruit1 fruit2
1) "apples"
```

### ▶ To remove any matching values in one set from another by finding the difference, use **SDIFF**:

```
127.0.0.1:6379> SDIFF fruit1 fruit2
1) "pears"
```

### ▶ To build a union of sets with values from both (since it's a set, any duplicates are dropped), use **SUNION**:

```
127.0.0.1:6379> SUNION fruit1 fruit2
1) "apples"
2) "oranges"
3) "pears"
```

20

20

## Data Types – Sorted Sets

### ▶ Sorted Sets

- Redis Sorted Sets are similar to Redis Sets, non-repeating collections of Strings. The difference is, every member of a Sorted Set is associated with a score, that is used in order to take the sorted set ordered, from the smallest to the greatest score. While members are unique, the scores may be repeated.

### ▶ Sorted Sets commands:

- ZADD key score1 member1 [score2 member2]** - Adds one or more members to a sorted set, or updates its score, if it already exists
- ZCARD key** - Gets the number of members in a sorted set
- ZRANGE key start stop [WITHSCORES]** - Returns a range of members in a sorted set, by index
- Full list sets commands see here:  
[https://redis.io/commands#sorted\\_set](https://redis.io/commands#sorted_set)
- Example:

```
127.0.0.1:6379> ZADD tutorials 1 redis
(integer) 1
127.0.0.1:6379> ZADD tutorials 2 mongodb
(integer) 1
127.0.0.1:6379> ZADD tutorials 3 mysql
(integer) 1
127.0.0.1:6379> ZADD tutorials 4 mysql
(integer) 0
127.0.0.1:6379> ZRANGE tutorials 0 10 WITHSCORES
1) "redis"
2) "1"
3) "mongodb"
4) "2"
5) "mysql"
6) "4"
```

21

## Questions



- ▶ What is **Redis-cli**?
- ▶ How can you save multiple values under one key in Redis?
- ▶ How to get value from Redis database?
- ▶ How will you delete a key from Redis?

22

22

## Part 3

23

23

## Transactions



- ▶ Redis transactions allow the execution of a group of commands in a single step.
- ▶ Redis transactions have the following properties:
  - All commands in a transaction are sequentially executed as a single isolated operation.
  - It is not possible that a request issued by another client is served in the middle of the execution of a Redis transaction.
- ▶ Redis transaction are also atomic.
  - Atomic means either all of the commands or none are processed.

24

24

# Transactions


- ▶ In Redis transactions are wrapped in **MULTI** block atomic commands.
  - Wrapping several operations in a single block will complete either successfully or not at all.
- ▶ When using **MULTI**, the commands aren't actually executed when we define them. Instead, they are queued and then executed in sequence.
- ▶ Similar to **ROLLBACK** in SQL, you can stop a transaction with the **DISCARD** command, which will clear the transaction queue.
  - Unlike **ROLLBACK**, it won't revert the database; it will simply not run the transaction at all.
  - The effect is identical, although the underlying concept is a different mechanism (transaction rollback vs. operation cancellation).
- ▶ Example:

```
127.0.0.1:6379> MULTI
OK
127.0.0.1:6379> SET count 2
QUEUED
127.0.0.1:6379> SET student Patel
QUEUED
127.0.0.1:6379> INCR count
QUEUED
127.0.0.1:6379> EXEC
1) OK
2) (integer) 3
```

25

25

# Namespaces



- ▶ Very often there is a need to separate keys by a namespace.
- ▶ In Redis a namespace is called a database and is keyed by number.
  - So far, we've always interacted with the default namespace 0 (also known as database 0).
  - To switch to another namespace you use **SELECT** command
  - You can move keys from one namespace to another with **MOVE** command.

```
127.0.0.1:6379> SET greeting hello
OK
127.0.0.1:6379> GET greeting
"hello"
127.0.0.1:6379> SELECT 1
OK
127.0.0.1:6379> GET greeting
(nil)
127.0.0.1:6379> SET greeting "ciao"
OK
127.0.0.1:6379> GET greeting
"ciao"
```

```
127.0.0.1:6379> SELECT 0
OK
127.0.0.1:6379> GET greeting
"hello"
127.0.0.1:6379> SELECT 1
OK
127.0.0.1:6379> GET greeting
"ciao"
```

26

26

## Expiry



- ▶ A common use case for a key-value system like Redis is as a fast-access cache for data that's more expensive to retrieve or compute.
- ▶ Expiration helps keep the total key set from growing unbounded, by tasking Redis to delete a key value after a certain time has passed.
- ▶ Marking a key for expiration requires the **EXPIRE** command, an existing key, and a time to live in seconds.
  - Redis also provides a shortcut command called **SETEX** to set the keys with expiration.

Example:

```
127.0.0.1:6379> SET ice "I'm melting..."
OK
127.0.0.1:6379> EXPIRE ice 10      //setting 10 sec expiry
(integer) 1

//check immediately
127.0.0.1:6379> EXISTS ice
(integer) 1

//check after 10sec
127.0.0.1:6379> EXISTS ice
(integer) 0
```

27

## Persistence

- ▶ Redis is very fast, but not very durable
- ▶ However, it provides a few persistence options:
  - **Snapshotting:**
    - Point-in-time snapshots of your dataset are created at specified intervals.
  - **The Append Only File:**
    - In that case Redis logs every write operation received by the server that will be played again at server startup, reconstructing the original dataset.
    - Commands are logged using the same format as the Redis protocol itself, in an append-only fashion. Redis is able to rewrite the log in the background when it gets too big.
  - It is possible to **combine both AOF and snapshots** in the same instance.
    - In this case, when Redis restarts the AOF file will be used to reconstruct the original dataset since it is guaranteed to be the most complete.
  - You can **disable persistence completely**, if you want your data to just exist as long as the server is running.

28

28

## Persistence – continued

### ► Snapshotting:

- By default Redis saves snapshots of the dataset on disk, in a binary file called **dump.rdb**.
- You can configure Redis to have it save the dataset every **N seconds** if there are at least **M changes** in the dataset,
- Or you can manually call the **SAVE** or **BGSAVE** commands.
- Example: `127.0.0.1:6379> SAVE 60 1000`
  - this will make Redis automatically dump the dataset to disk **every 60 seconds** if at least **1000 keys changed**

### ► Append-only file:

- The append-only file is an alternative, fully-durable strategy for Redis. It became available in version 1.1.
  - You can turn on the AOF in your configuration file using: **appendonly yes**
  - In that case every time Redis receives a command that changes the dataset (e.g. **SET**) it will append it to the AOF.
  - When you restart Redis it will re-play the AOF to rebuild the state.

More Info here <https://redis.io/topics/persistence>

29

29

## Questions



- What is the purpose of a Namespace in Redis?
- What is the difference between SET and MSET commands in Redis?
- What is the purpose of SELECT command in Redis?
- How to create a key that expires after 300 seconds and hold a string value?

30

30

## Further reading

- ▶ Carlson J. Redis in Action. Manning Publications, 2013
- ▶ Kreibich J. Redis: The Definitive Guide: Data modeling, caching, and messaging, O'Reilly Media, 2015
- ▶ Macedo T., Oliveira F. Redis Cookbook: Practical Techniques for Fast Data Manipulation, O'Reilly Media, 2011
- ▶ Redis Documentation <https://redis.io/>



31

31

## Essentials

- ✓ Introduced Redis data store
- ✓ Discussed CRUD commands
- ✓ Introduced different data types
- ✓ Learned to work with complex data types
- ✓ Discussed Redis transactions
- ✓ Discussed the use of the expiration

32

32